

CS2102 Lecture 3

SQL (Part 1)

Structured Query Language

- Designed by D. Chamberlin & R. Boyce (IBM Research) in 1974
 - Original name: SEQUEL (Structured English QUERy Language)
- SQL is not a general-purpose language but a **domain-specific language (DSL)**
 - Designed for computations on relations
- Unlike relational algebra which is a procedural language, SQL is a **declarative language**
 - Focuses on *what* to compute, not on *how* to compute

Using SQL

- **Directly write SQL statements**

- Command line interface

- PostgreSQL's **psql**

- <https://www.postgresql.org/docs/current/static/app-psql.html>

- Graphical user interface

- E.g., PostgreSQL's pgAdmin <https://www.pgadmin.org/>

- **Include SQL in application programs**

- Statement-Level Interface (SLI)

- Embedded SQL
 - Dynamic SQL

- Call-Level Interface (CLI)

- JDBC (Java DataBase Connectivity)
 - ODBC (Open DataBase Connectivity)

SQL

- Current ANSI/ISO standard for SQL is called SQL:2019
 - Different DBMSs may have minor variations in SQL syntax
- SQL consists of two main parts: DDL & DML
- **Data Definition Language (DDL):**
create/delete/modify schemas
- **Data Manipulation Language (DML):** ask queries,
insert/delete/modify data

Create/Drop Table

-- Comments in SQL are preceded
-- by two hyphens

```
create table Students (  
    studentId      integer,  
    name           varchar(100),  
    birthDate      date  
);
```

/* SQL also supports
C-style comments */

```
drop table Students;  
drop table if exists Students;  
drop table if exists Students cascade;
```

Data Types

- Examples of built-in data types:

boolean	false/true (null represents unknown)
integer	signed four-byte integer
float8	double-precision floating point number (8 bytes)
numeric	arbitrary precision floating point number
numeric(p,s)	maximum total of p digits with maximum of s digits in fractional part
char(n)	fixed-length string consisting of n characters
varchar(n)	variable-length string up to n characters
text	variable-length character string
date	calendar date (year, month, day)
timestamp	date and time

- Other data types: enumerated data type, array, XML, JSON, user-defined data type, etc.
- Reference: <https://www.postgresql.org/docs/current/datatype.html>

Null Values

- Result of a **comparison operation** involving *null* value is *unknown*
- Result of an **arithmetic operation** involving *null* value is *null*
- Example: Assume that the value of *x* is *null*
 - $x < 100$ evaluates to *unknown*
 - $x = null$ evaluates to *unknown*
 - $x <> null$ evaluates to *unknown*
 - $x + 20$ evaluates to *null*

Null Values (cont.)

- SQL uses a **three-valued logic system**: *true*, *false*, & *unknown*

x	y	x AND y	x OR y	NOT x
FALSE	FALSE	FALSE	FALSE	TRUE
FALSE	UNKNOWN	FALSE	UNKNOWN	
FALSE	TRUE	FALSE	TRUE	
UNKNOWN	FALSE	FALSE	UNKNOWN	UNKNOWN
UNKNOWN	UNKNOWN	UNKNOWN	UNKNOWN	
UNKNOWN	TRUE	UNKNOWN	TRUE	
TRUE	FALSE	FALSE	TRUE	FALSE
TRUE	UNKNOWN	UNKNOWN	TRUE	
TRUE	TRUE	TRUE	TRUE	

IS NULL Comparison Predicate

- How to check if a value is equal to *null*?
- Use the **IS NULL** comparison predicate
 - *x IS NULL* evaluates to *true* if *x* is a null value
 - *x IS NULL* evaluates to *false* if *x* is a non-null value
- “*x IS NOT NULL*” is equivalent to “**NOT (x IS NULL)**”

IS DISTINCT FROM Comparison Predicate

- How to treat *null* values as ordinary values for comparison?
- Use the **IS DISTINCT FROM** comparison predicate
- The comparison “**x IS DISTINCT FROM y**”
 - is equivalent to “ $x \neq y$ ” if both *x* & *y* are non-null values
 - evaluates to *false* if both the values are *null*
 - evaluates to *true* if only one of the values is *null*
- “**x IS NOT DISTINCT FROM y**” is equivalent to “**NOT (x IS DISTINCT FROM y)**”

Constraints in Data Definitions

- **Constraint Types**
 - Not-null constraints
 - Unique constraints
 - Primary key constraints
 - Foreign key constraints
 - General constraints
- **Constraint Specifications**
 - Column constraints
 - Table constraints
 - Assertions
- A constraint is **violated** if it evaluates to *false*

Not-Null Constraints

```
create table Students (  
    studentId      integer,  
    name           varchar(100) not null,  
    birthDate      date  
);
```

The not-null constraint is violated if there exists some record $x \in Students$ where “**x.name IS NOT NULL**” evaluates to *false*

Unique Constraints

```
create table Students (  
    studentId      integer unique,  
    name           varchar(100),  
    birthDate      date  
);
```

The unique constraint is violated if there exists two records $x, y \in Students$ where “ $x.studentId \neq y.studentId$ ” evaluates to *false*

studentId	name	birthDate
100	Alice	1999-12-25
200	Bob	2000-04-01
200	Carol	2001-11-28

studentId	name	birthDate
100	Alice	1999-12-25
null	Bob	2000-04-01
null	Carol	2001-11-28

Unique Constraints (cont.)

```
create table Census (  
    city      varchar(50),  
    state     char(2),  
    population integer,  
    unique (city,state)  
);
```

city	state	population
New York	NY	8537673
null	CA	3976322
Chicago	IL	2704958
Houston	TX	2303482
null	CA	1406630

The unique constraint is violated if there exists two records $x, y \in \text{Census}$ where
“(x.city \neq y.city) or (x.state \neq y.state)”
evaluates to *false*

Primary Key Constraints

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100),  
    birthDate      date  
);
```

```
create table Students (  
    studentId      integer unique not null,  
    name           varchar(100),  
    birthDate      date  
);
```

Primary Key Constraints (cont.)

```
create table Enrolls (  
    sid      integer,  
    cid      integer,  
    grade    char(2),  
    primary key (sid, cid)  
);
```


Foreign Key Constraints

```
create table Students (  
  studentId      integer,  
  name           varchar(100),  
  birthDate      date,  
  primary key (studentId));
```

```
create table Courses (  
  courseId       integer,  
  name           varchar(80),  
  credits         integer,  
  primary key (courseId));
```

```
create table Enrolls (  
  sid            integer references Students (studentId),  
  cid            integer,  
  grade          char(2),  
  primary key (sid, cid),  
  foreign key (cid) references Courses (courseId));
```

Foreign Key Constraints (cont.)

```
create table Rooms (  
    level            integer,  
    number           integer,  
    capacity         integer,  
    primary key (level, number)  
);
```

```
create table Events (  
    eid             integer,  
    ename           varchar(100),  
    date            date,  
    level           integer,  
    number          integer,  
    primary key (eid),  
    foreign key (level, number)  
        references Rooms (level, number)  
);
```

Rooms		
level	number	capacity
1	1	30
1	2	100
2	1	60
2	2	20
2	3	70

Events				
eid	ename	date	level	number
101	CS Seminar	2018-01-02	1	2
102	PhD Defense	2018-03-10	2	3
103	Job Fair	2018-03-14	2	null
104	CS Interviews	2018-06-24	null	null
105	CS Meeting	2018-06-20	null	9

Foreign Key Constraints (cont.)

```
create table Rooms (  
    level          integer,  
    number         integer,  
    capacity       integer,  
    primary key (level, number)  
);
```

Rooms		
level	number	capacity
1	1	30
1	2	100
2	1	60
2	2	20
2	3	70

```
create table Events (  
    eid            integer,  
    ename         varchar(100),  
    date          date,  
    level         integer,  
    number        integer,  
    primary key (eid),  
    foreign key  (level, number)  
                references Rooms (level, number)  
                match full  
);
```

Events				
eid	ename	date	level	number
101	CS Seminar	2018-01-02	1	2
102	PhD Defense	2018-03-10	2	3
104	CS Interviews	2018-06-24	null	null

Foreign Key Constraints (cont.)

```
create table Rooms (  
  level      integer,  
  number     integer,  
  capacity   integer,  
  primary key (level, number),  
);
```

```
create table Events (  
  eid        integer,  
  ename      varchar(100),  
  date       date,  
  level      integer,  
  number     integer,  
  primary key (eid),  
  foreign key (level, number)  
    references Rooms (level, number)  
);
```

```
create table Rooms (  
  level      integer,  
  number     integer,  
  capacity   integer,  
  primary key (level, number),  
);
```

```
create table Events (  
  eid        integer,  
  ename      varchar(100),  
  date       date,  
  level      integer,  
  number     integer,  
  primary key (eid),  
  foreign key (level, number)  
    references Rooms  
);
```

Foreign Key Constraints (cont.)

```
create table Rooms (  
    rid          integer,  
    level        integer,  
    number       integer,  
    capacity     integer,  
    primary key (rid),  
    unique (level, number)  
);  
  
create table Events (  
    eid          integer,  
    ename        varchar(100),  
    date         date,  
    level        integer,  
    number       integer,  
    primary key (eid),  
    foreign key (level, number) references Rooms (level, number)  
);
```

Check Constraints

```
create table Lectures (  
  cname      char(5),  
  pname      varchar(50) not null,  
  day        smallint  
             check (day in (1,2,3,4,5)),  
  hour       smallint  
             check ((hour >= 8) and (hour <= 17)),  
  primary key (cname,day,hour),  
  unique (pname,day,hour)  
);
```

Check Constraints (cont.)

```
create table Lectures (  
    cname      char(5),  
    pname      varchar(50) not null,  
    day        smallint,  
    hour       smallint,  
    primary key (cname,day,hour),  
    unique (pname,day,hour),  
    check (  
        (day in (1,2,3,4,5) and hour >= 8 and hour <= 17)  
        or  
        (day = 6 and hour in (9,10,11))  
    )  
);
```

Constraint Names

```
create table Lectures (  
  cname      char(5),  
  pname      varchar(50)  
             constraint lectures_pname check (pname is not null),  
  day        smallint  
             constraint lectures_day check (day in (1,2,3,4,5)),  
  hour       smallint  
             constraint lectures_hour check (hour >= 8 and hour <= 17),  
  constraint lectures_pri_key primary key (cname,day,hour),  
  constraint lectures_cand_key unique (pname,day,hour)  
);
```


Assertions

- SQL's **create assertion** statement supports the specification of general constraints
 - Example: Every course is enrolled by at least one student
- However, the create assertion statement is not implemented in many DBMSs
- Instead, general constraints are enforced using **triggers**

Database Modifications: Insert

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    dept           varchar(20)  
);
```

```
insert into Students  
values      (12345, 'Alice', '1999-12-25', 'Maths');
```

```
insert into Students (name, studentId)  
values      ('Bob', 67890), ('Carol', 11122);
```

studentId	name	birthDate	dept
12345	Alice	1999-12-25	Maths
67890	Bob	null	null
11122	Carol	null	null

Database Modifications: Insert (cont.)

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    dept           varchar(20) default 'CS'  
);
```

```
insert into Students  
values      (12345, 'Alice', '1999-12-25', 'Maths');
```

```
insert into Students (name, studentId)  
values      ('Bob', 67890), ('Carol', 11122);
```

studentId	name	birthDate	dept
12345	Alice	1999-12-25	Maths
67890	Bob	null	CS
11122	Carol	null	CS

Database Modifications: Delete

```
create table Students (  
    studentId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    dept           varchar(20) default 'CS'  
);
```

```
-- Remove all students  
delete from Students;
```

```
-- Remove all students from Maths department  
delete from Students  
where      dept = 'Maths';
```

Database Modifications: Update

```
create table Accounts (  
    accountId      integer primary key,  
    name           varchar(100) not null,  
    birthDate      date,  
    balance        numeric (10,2) default 0.00  
);
```

-- Add 2% interest to all accounts

```
update Accounts  
set     balance = balance * 1.02;
```

-- Add \$500 to account 12345 & change name to 'Alice'

```
update   Accounts  
set      balance = balance + 500,  
          name = 'Alice'  
where    accountId = 12345;
```

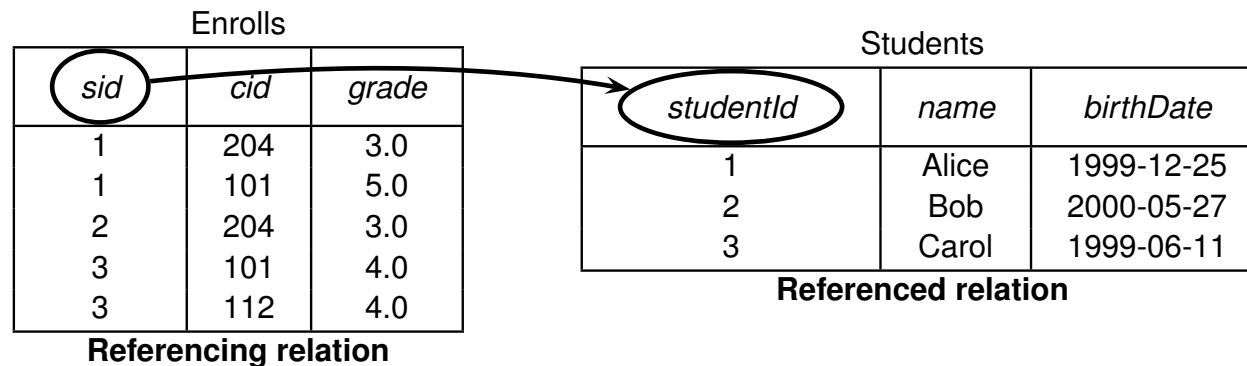
How to retrieve all the records from a table?

```
select * from Customers;
```

Customers

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

Foreign Key Constraint Violations



- Deletion/update of a referenced tuple could violate a foreign key constraint
- Specify action to deal with violation as part of a foreign key constraint declaration:
FOREIGN KEY ... REFERENCES ... ON DELETE/UPDATE action
- Default action is “NO ACTION”
 - The delete/update statement that causes violation will be rejected

Foreign Key Constraint Violations (cont.)

- **NO ACTION**: rejects delete/update if it violates constraint (default option)
- **RESTRICT**: similar to *NO ACTION* except that constraint checking can't be deferred
- **CASCADE**: propagates delete/update to referencing tuples
- **SET DEFAULT**: updates foreign keys of referencing tuples to some default value
- **SET NULL**: updates foreign keys of referencing tuples to *null* value

Foreign Key Constraint Violations (cont.)

```
create table Students (  
  studentId      integer,  
  name           varchar(100),  
  birthDate      date,  
  primary key (studentId));
```

```
create table Courses (  
  courseId       integer,  
  name           varchar(80),  
  credits         integer,  
  primary key (courseId));
```

```
create table Enrolls (  
  sid integer,  cid integer,  grade char(2),  
  primary key (sid, cid),  
  foreign key (sid) references Students  
    on delete cascade  
    on update cascade,  
  foreign key (cid) references Courses);
```

Foreign Key Constraint Violations (cont.)

foreign key (sid) references Students on delete cascade

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>
1	Alice	1999-12-25
2	Bob	2000-05-27
3	Carol	1999-06-11

delete from
Students
where studentId = 1;

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>
2	Bob	2000-05-27
3	Carol	1999-06-11

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
1	204	3.0
1	101	5.0
2	204	3.0
3	101	4.0
3	112	4.0

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
2	204	3.0
3	101	4.0
3	112	4.0

Foreign Key Constraint Violations (cont.)

foreign key (sid) references Students on update cascade

<i>studentId</i>	<i>name</i>	<i>birthDate</i>
1	Alice	1999-12-25
2	Bob	2000-05-27
3	Carol	1999-06-11

update Students
set studentId = 4
where studentId = 1;

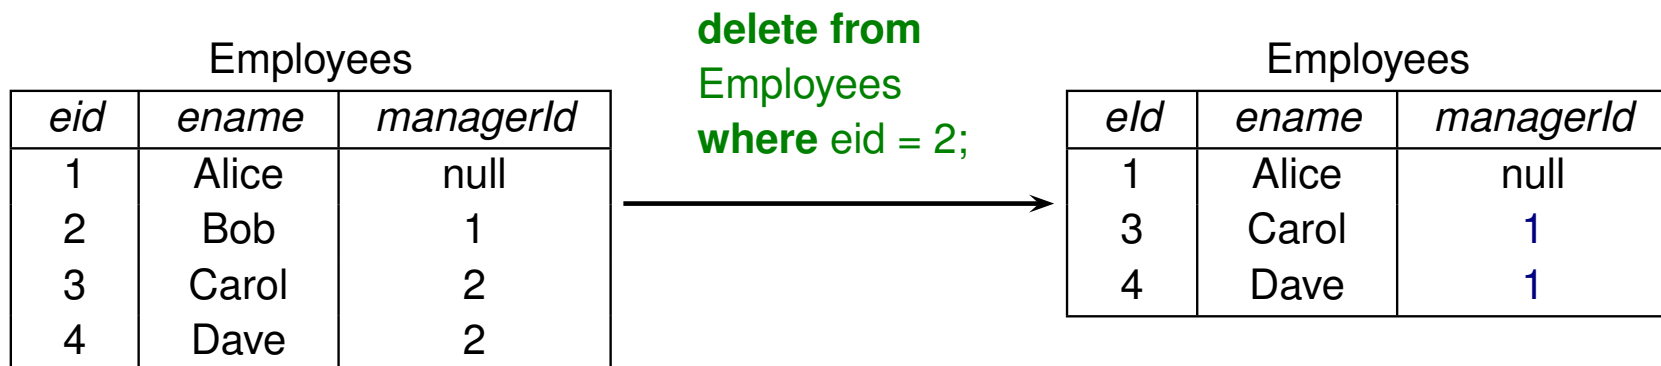
<i>studentId</i>	<i>name</i>	<i>birthDate</i>
4	Alice	1999-12-25
2	Bob	2000-05-27
3	Carol	1999-06-11

<i>sid</i>	<i>cid</i>	<i>grade</i>
1	204	3.0
1	101	5.0
2	204	3.0
3	101	4.0
3	112	4.0

<i>sid</i>	<i>cid</i>	<i>grade</i>
4	204	3.0
4	101	5.0
2	204	3.0
3	101	4.0
3	112	4.0

Foreign Key Constraint Violations (cont.)

```
create table Employees (  
    eid          integer primary key,  
    ename        varchar(100),  
    managerId    integer default 1,  
    foreign key (managerId) references Employees  
                on delete set default  
);
```



Transactions

- A **transaction** consists of one or more update/retrieval operations (i.e., SQL statements)
- Abstraction for representing a logical unit of work
- The **begin** command starts a new transaction
- Each transaction must end with either a **commit** or **rollback** command

```
begin;  
update Accounts  
set balance = balance + 1000  
where accountId = 2;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 1;  
commit;
```

Transactions (cont.)

```
begin;  
update Accounts  
set balance = balance + 1000  
where accountId = 2;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 1;  
commit;
```

vs.

```
update Accounts  
set balance = balance + 1000  
where accountId = 2;  
  
update Accounts  
set balance = balance - 1000  
where accountId = 1;
```

Transactions: ACID Properties

- **Atomicity**: Either all the effects of a transaction are reflected in the database or none are
- **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
- **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
- **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

Deferrable Constraints

- By default, constraints are checked **immediately** at the end of SQL statement execution
 - A violation will cause the statement to be rolledback

```
create table Employees (  
    eid          integer primary key,  
    ename        varchar(100),  
    managerId    integer,  
    foreign key (managerId) references Employees  
);
```

```
insert into Employees values (1, 'Alice', null), (2, 'Bob', 1), (3, 'Carol', 2);
```

```
delete from Employees where eid = 2;
```

```
update Employees set managerId = 1 where eid = 3;
```


Deferrable Constraints (cont.)

- The checking could also be **deferred** for some constraints to the **end of transaction** execution
 - A violation will cause the transaction to be aborted
 - Deferrable constraints:
 - unique
 - primary key
 - foreign key
- Deferrable constraints are specified using
 - either **deferrable initially deferred**
 - or **deferrable initially immediate**
- The checking of deferrable constraints can be changed using **set constraints** statement

Deferrable Constraints (cont.)

```
create table Employees (  
    eid          integer primary key,  
    ename        varchar(100),  
    managerId    integer,  
    foreign key (managerId) references Employees  
                deferrable initially deferred  
);
```

```
insert into Employees values (1, 'Alice', null), (2, 'Bob', 1), (3, 'Carol', 2);
```

```
begin;
```

```
delete from Employees where eid = 2;
```

```
update Employees set managerId = 1 where eid = 3;
```

```
commit;
```

Deferrable Constraints (cont.)

```
create table Employees (  
    eid            integer primary key,  
    ename          varchar(100),  
    managerId      integer,  
    constraint employees_fkey foreign key (managerId) references Employees  
        deferrable initially immediate  
);
```

```
insert into Employees values (1, 'Alice', null), (2, 'Bob', 1), (3, 'Carol', 2);
```

```
begin;
```

```
set constraints employees_fkey deferred;
```

```
delete from Employees where eid = 2;
```

```
update Employees set managerId = 1 where eid = 3;
```

```
commit;
```

Modifying Schema

- Add/remove/modify columns
 - **alter table** Students **alter column** dept **drop default**;
 - **alter table** Students **drop column** dept;
 - **alter table** Students **add column** faculty **varchar**(20);
 - etc.
- Add/remove constraints
 - **alter table** Students **add constraint** fk_grade **foreign key** (grade) **references** Grades;
 - etc.
- etc.
- Reference:

<https://www.postgresql.org/docs/current/sql-altertable.html>

Summary

- SQL is the standard query language for relational DBMS
- Column/Table Constraints
 - not null
 - unique
 - primary key
 - foreign key
 - check
- Transactions & deferrable constraints
- SQL Reference:
<https://www.postgresql.org/docs/current/index.html>