

CS2102 Lecture 6

SQL (Part 3)

Aggregate Functions

- Aggregate function computes a single value from a set of tuples
- **Example:** Find the minimum, maximum, and average prices of pizzas sold by Corleone Corner

```
select  min (price), max (price), avg (price)
from    Sells
where   rname = 'Corleone Corner'
```

Sells		
rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

min	max	avg
19	25	22.666666666666667

Aggregate Functions (cont.)

Basic aggregate functions: min, max, avg, sum, count

Query	Meaning
select min(A) from R	Minimum non-null value in A
select max(A) from R	Maximum non-null value in A
select avg(A) from R	Average of non-null values in A
select sum(A) from R	Sum of non-null values in A
select count(A) from R	Count number of non-null values in A
select count(*) from R	Count number of rows in R
select avg(distinct A) from R	Average of distinct non-null values in A
select sum(distinct A) from R	Sum of distinct non-null values in A
select count(distinct A) from R	Count number of distinct non-null values in A

For more aggregate functions, refer to

<https://www.postgresql.org/docs/current/functions-aggregate.html>

Aggregate Functions (cont.)

- Let R be an empty relation
- Let S be a relation with cardinality = n where all values of attribute A are null values

Query	Result
select min(A) from R	null
select max(A) from R	null
select avg(A) from R	null
select sum(A) from R	null
select count(A) from R	0
select count(*) from R	0

Query	Result
select min(A) from S	null
select max(A) from S	null
select avg(A) from S	null
select sum(A) from S	null
select count(A) from S	0
select count(*) from S	n

Usage of Aggregate Functions

- Aggregate functions can be used in different parts of SQL queries:
 - SELECT clause
 - HAVING clause (to be discussed later)
 - ORDER BY clause (to be discussed later)

Usage of Aggregate Functions (cont.)

Find the number of items ordered and the maximum order cost for an item

Orders

item	price	qty
A	2.50	100
B	4.00	100
C	7.50	100

count	max
3	750.00

```
select count(*), max(price * qty)
from Orders;
```

Usage of Aggregate Functions (cont.)

Find the most expensive pizzas and the restaurants that sell them (at the most expensive price)

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	25
Pizza King	Diavola	17
Pizza King	Hawaiian	21

pizza	rname
Hawaiian	Corleone Corner
Marinara	Mamma's Place

```
select  pizza, rname
from    Sells
where   price = (select max(price) from Sells);
```

GROUP BY Clause

For each restaurant that sells some pizza, find the minimum and maximum prices of its pizzas

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	min	max
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

GROUP BY Clause (cont.)

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	min	max
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

Conceptual processing steps:

1. Partition the tuples in Sells into groups based on rname
2. Compute min(price) and max(price) for each group
3. Output one tuple for each group

GROUP BY Clause (cont.)

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	min	max
Corleone Corner	19	25
Gambino Oven	16	16
Lorenzo Tavern	23	23
Mamma's Place	22	22
Pizza King	17	21

```
select    rname, min(price), max(price)
from      Sells
group by rname;
```

GROUP BY Clause (cont.)

Find the number of students for each (dept,year) combination. Show the output in ascending order of (dept,year).

Students

studentId	name	year	dept
12345	Alice	1	Maths
60031	George	1	Maths
18763	Fred	3	Maths
11123	Carol	4	Maths
67890	Bob	2	CS
87012	Hugh	2	CS
20135	Eve	3	CS
20135	Dave	4	CS
96410	Ivy	4	CS

dept	year	num
CS	2	2
CS	3	1
CS	4	2
Maths	1	2
Maths	3	1
Maths	4	1

```
select    dept, year, count(*) as num
from      Students
group by  dept, year;
order by  dept, year;
```

GROUP BY Clause (cont.)

Show all restaurants in descending order of their average pizza price. Exclude restaurants that do not sell any pizza.

```
select    rname
from      Sells
group by  rname
order by  avg(price) desc;
```

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname
Lorenzo Tavern
Corleone Corner
Mamma's Place
Pizza King
Gambino Oven

GROUP BY Clause (cont.)

For each restaurant that sells some pizza, find its average pizza price. Show the restaurants in descending order of their average pizza price.

```
select    rname, avg(price) as avg_price
from      Sells
group by  rname
order by  avg_price desc;
```

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname	avg_price
Lorenzo Tavern	23.0000000000000000
Corleone Corner	22.6666666666666667
Mamma's Place	22.0000000000000000
Pizza King	19.0000000000000000
Gambino Oven	16.0000000000000000

GROUP BY Clause: Properties

- In a query with “GROUP BY a_1, a_2, \dots, a_n ”, two tuples t & t' belong to the same group if the following expression evaluates to true:

$(t.a_1 \text{ IS NOT DISTINCT FROM } t'.a_1) \text{ AND } \dots \text{ AND } (t.a_n \text{ IS NOT DISTINCT FROM } t'.a_n)$

- **Example:** Four groups in R if R is grouped by $\{A, C\}$

R		
A	B	C
null	4	19
null	21	19
6	1	null
6	20	null
20	2	10
1	1	2
1	18	2

GROUP BY Clause: Properties (cont.)

These queries are invalid!

Q1: select year, **count**(*)
 from Students
 group by dept;

Q2: select dept, **count**(*)
 from Students
 group by dept;
 order by year;

Students

studentId	name	year	dept
12345	Alice	1	Maths
11123	Carol	4	Maths
18763	Fred	3	Maths
60031	George	1	Maths
67890	Bob	2	CS
20135	Dave	4	CS
20135	Eve	3	CS
87012	Hugh	2	CS
96410	Ivy	4	CS

GROUP BY Clause: Properties (cont.)

rname	area	rname	pizza	price
Corleone Corner	North	Corleone Corner	Diavola	24
Corleone Corner	North	Corleone Corner	Hawaiian	25
Corleone Corner	North	Corleone Corner	Margherita	19
Gambino Oven	Central	Gambino Oven	Siciliana	16
Lorenzo Tavern	Central	Lorenzo Tavern	Funghi	23
Mamma's Place	South	Mamma's Place	Marinara	22
Pizza King	East	Pizza King	Diavola	17
Pizza King	East	Pizza King	Hawaiian	21

rname	area	avg_price
Corleone Corner	North	22.67
Gambino Oven	Central	16.00
Lorenzo Tavern	Central	23.00
Mamma's Place	South	22.00
Pizza King	East	19.00

```
select      R.rname, R.area, round(avg(S.price),2) as avg_price
from        Restaurants R, Sells S
where       S.rname = R.rname
group by    R.rname;
```

```
select      rname, area, round(avg(price),2) as avg_price
from        Restaurants natural join Sells
group by    rname;
```


GROUP BY Clause: Properties (cont.)

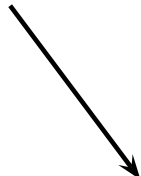
rname	area	rname	pizza	price
Corleone Corner	North	Corleone Corner	Diavola	24
Corleone Corner	North	Corleone Corner	Hawaiian	25
Corleone Corner	North	Corleone Corner	Margherita	19
Gambino Oven	Central	Gambino Oven	Siciliana	16
Lorenzo Tavern	Central	Lorenzo Tavern	Funghi	23
Mamma's Place	South	Mamma's Place	Marinara	22
Pizza King	East	Pizza King	Diavola	17
Pizza King	East	Pizza King	Hawaiian	21

rname	area	avg_price
Corleone Corner	North	22.67
Gambino Oven	Central	16.00
Lorenzo Tavern	Central	23.00
Mamma's Place	South	22.00
Pizza King	East	19.00

select S.rname, R.area, **round**(avg(S.price),2) **as** avg_price
from Restaurants R, Sells S
where S.rname = R.rname
group by S.rname;

select rname, area, **round**(avg(price),2) **as** avg_price
from Sells **natural join** Restaurants
group by rname;

GROUP BY Clause: Properties (cont.)

- Each output tuple corresponds to one group
 - For each column A in relation R that appears in the SELECT clause, one of the following conditions must hold:
 1. A appears in the GROUP BY clause,
 2. A appears in an aggregated expression in the SELECT clause (e.g., **min**(A)), or
 3. the primary ~~(or a candidate)~~ key of R appears in the GROUP BY clause
- 
Not supported in PostgreSQL
- For this module, we will follow PostgreSQL's more restrictive group-by clause properties

GROUP BY Clause: Properties (cont.)

The following query is valid in standard SQL but invalid in PostgreSQL

```
select    pname
from      Lectures
group by cname, day, hour;
```

- Primary key of Lectures: $\{pname, day, hour\}$
- Candidate key of Lectures: $\{cname, day, hour\}$

Lectures

cname	pname	day	hour
CS101	Alice	1	10
CS123	Alice	1	15
CS123	Alice	3	15
CS200	Bob	4	8
MA300	Bob	3	15

GROUP BY Clause: Properties (cont.)

- If an aggregate function appears in the SELECT clause and there is no GROUP BY clause, then the SELECT clause must not contain any column that is not in an aggregated expression
- **Example:** The following query is invalid!

```
select    rname, min(price), max(price)
from      Sells
```

Removing Duplicate Records with GROUP BY

**Q1: select distinct rname
from Sells;**

**Q2: select rname
from Sells
group by rname;**

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

rname
Corleone Corner
Gambino Oven
Lorenzo Tavern
Mamma's Place
Pizza King

Quiz

Q1: **select** rname
 from Sells
 order by price;

Q2: **select** **distinct** rname
 from Sells
 order by price;

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

Output of Q1

rname
Gambino Oven
Pizza King
Corleone Corner
Pizza King
Mamma's Place
Lorenzo Tavern
Corleone Corner
Corleone Corner

- Query Q1 is valid but query Q2 is invalid. Why?

HAVING Clause

Find restaurants that sell pizzas with an average selling price of at least \$22

Sells

rname	pizza	price
Corleone Corner	Diavola	24
Corleone Corner	Hawaiian	25
Corleone Corner	Margherita	19
Gambino Oven	Siciliana	16
Lorenzo Tavern	Funghi	23
Mamma's Place	Marinara	22
Pizza King	Diavola	17
Pizza King	Hawaiian	21

avg(price) = 22.67

avg(price) = 19

rname
Corleone Corner
Lorenzo Tavern
Mamma's Place

```
select    rname
from      Sells
group by  rname
having   avg(price) >= 22;
```

HAVING Clause (cont.)

Find restaurants located in the 'East' area that sell pizzas with an average selling price higher than the minimum selling price at Pizza King

```
select    rname
from      Sells
where      rname in (
              select rname
              from Restaurants
              where area = 'East'
            )
group by  rname
having     avg(price) >
            (select min(price)
             from    Sells
             where   rname = 'Pizza King');
```


HAVING Clause: Properties

- For each column A in relation R that appears in the HAVING clause, one of the following conditions must hold:
 - A appears in the GROUP BY clause,
 - A appears in an aggregated expression in the HAVING clause, or
 - the primary (or a candidate) key of R appears in the GROUP BY clause

Students

studentId	name	year	dept
12345	Alice	1	Maths
11123	Carol	4	Maths
18763	Fred	3	Maths
60031	George	1	Maths
67890	Bob	2	CS
20135	Dave	4	CS
20135	Eve	3	CS
87012	Hugh	2	CS
96410	Ivy	4	CS

This query is invalid!

```
select    dept, count(*)
from      Students
group by  dept
having    year = 3;
```

Quiz

Q1: **select** rname, **avg**(price)
from (Restaurants **natural join** Sells) RS
group by rname
having **avg**(price) > (
select **avg**(price)
from Sells **natural join** Restaurants
where area = RS.area
);

Q2: **select** rname, **avg**(price)
from Restaurants R **natural join** Sells
group by R.rname
having **avg**(price) > (
select **avg**(price)
from Sells **natural join** Restaurants
where area = R.area
);

Q3: **select** rname, **avg**(price)
from (Sells **natural join** Restaurants) SR
group by rname
having **avg**(price) > (
select **avg**(price)
from Sells **natural join** Restaurants
where area = SR.area
);

Q4: **select** rname, **avg**(price)
from Sells S **natural join** Restaurants R
group by S.rname
having **avg**(price) > (
select **avg**(price)
from Sells **natural join** Restaurants
where area = R.area
);

Group By Queries

Q1: **select** rname, **avg**(price)
from Sells
group by rname;

Q2: **select** rname
from Sells
group by rname;

Q3: **select** **avg**(price)
from Sells;

Q4: **select** **avg**(price)
from Sells;
having **count**(*) > 0;

Conceptual Evaluation of Queries

select	distinct select-list
from	from-list
where	where-condition
group by	groupby-list
having	having-condition
order by	orderby-list
offset	offset-specification
limit	limit-specification

1. Compute the cross-product of the tables in **from-list**
2. Select the tuples in the cross-product that evaluate to *true* for the **where-condition**
3. Partition the selected tuples into groups using the **groupby-list**
4. Select the groups that evaluate to *true* for the **having-condition** condition
5. For each selected group, generate an output tuple by selecting/computing the attributes/expressions that appear in the **select-list**
6. Remove any duplicate output tuples
7. Sort the output tuples based on the **orderby-list**
8. Remove the appropriate output tuples based on the **offset-specification** & **limit-specification**

Conceptual Evaluation of Queries (cont.)

```
select    rname, price as x  
from      Sells  
where     x > 30;
```

- The reference to x in the WHERE clause is invalid

Conceptual Evaluation of Queries (cont.)

Q1: **select** rname, **avg**(price) **as** avg_price
from Sells
group by rname
having **avg**(price) > 30;

Q2: **select** rname, **avg**(price) **as** avg_price
from Sells
group by rname
having avg_price > 30;

- Query Q1 is valid but query Q2 is invalid

Usage of Aggregate Functions

Q1: **select** rname, **min**(price) **as** min_price
from Sells
group by rname
having **avg**(price) > 30;
order by **sum**(price);

Q2: **select** *
from Sells
where price = **max**(price);

Q3: **select** *
from Sells
where price = (**select max**(price) **from** Sells);

- Queries Q1 & Q3 are valid but query Q2 is invalid

Common Table Expressions (CTEs)

Find restaurants where the average selling price of its pizzas is higher than the average selling price of pizzas in that restaurant's area.

```
with rname_avgprice as (  
    select rname, avg(price) as avg_price  
    from Sells  
    group by rname  
)  
area_avgprice as (  
    select area, avg(price) as avg_price  
    from Sells natural join Restaurants  
    group by area  
)  
select rname  
from rname_avgprice R  
where avg_price > (  
    select avg_price  
    from area_avgprice  
    where area = (select area from Restaurants where rname = R.rname)  
);
```


Common Table Expressions (CTEs)

with

R1 **as** (Q1),

R2 **as** (Q2),

...

Rn **as** (Qn)

select/insert/update/delete statement *S*;

- Each R_i is the name of a temporary relation defined by a query Q_i .
- Each R_i can reference to any of the preceding relations R_j , $j < i$
- S is a SQL statement that references R_n & possibly R_1, R_2, \dots
- CTEs can be used for writing **recursive queries** (not covered)

Views

- A **view** defines a virtual relation that can be used for querying
- **Example:** Consider the following database schema:
Courses (courseId, cname, credits, profId, lectureTime, quota)
Profs (profId, pname, officeRoom, contactNum)
Students (studentId, sname, email, birthDate)
Enrollment (courseId, numUGrad, numPGrad, numExchange, numAudit)

create view CourseInfo **as**

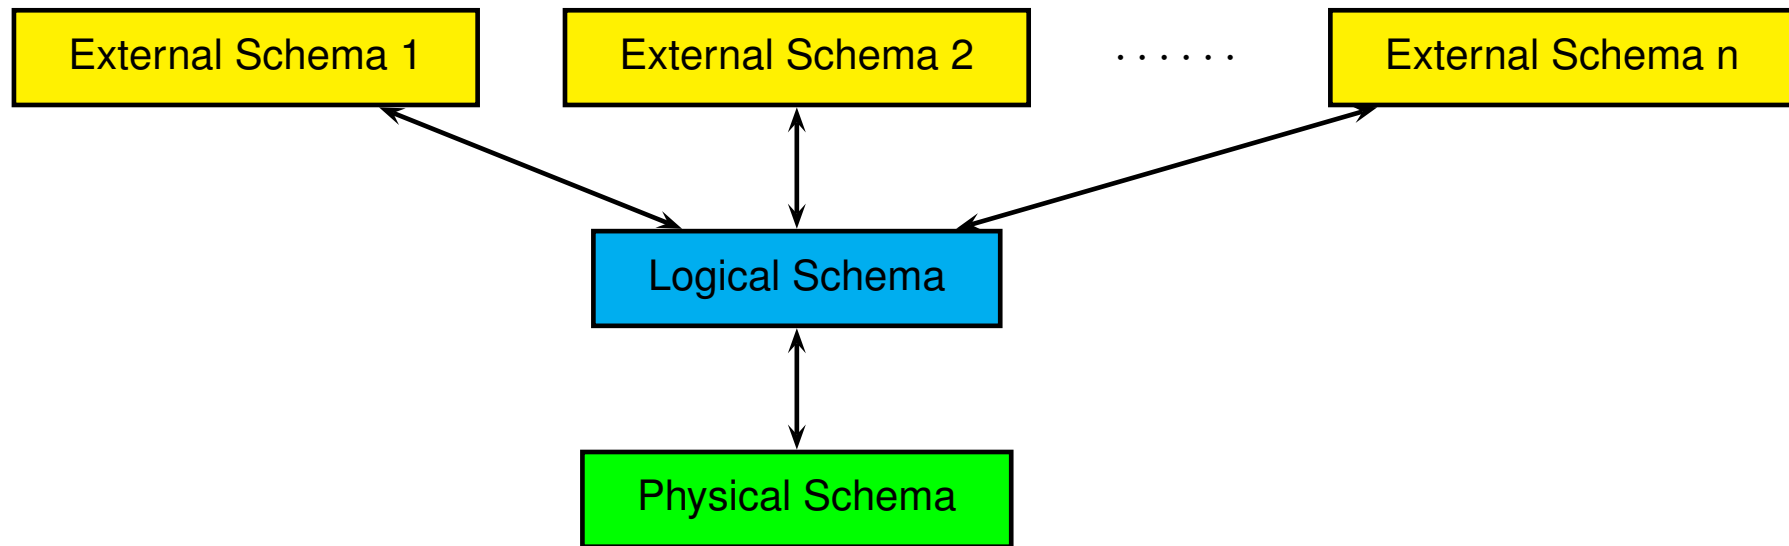
select cname, pname, lectureTime,
numUGrad+numPGrad+numExchange+numAudit **as** numEnrolled
from Courses **natural join** Profs **natural join** Enrollment;

Views (cont.)

```
create view CourseInfo as  
  select cname, pname, lectureTime,  
    numUGrad+numPGrad+numExchange+numAudit as numEnrolled  
from   Courses natural join Profs natural join Enrollment;
```

```
create view CourseInfo (cname, pname, lectureTime, numEnrolled) as  
  select cname, pname, lectureTime,  
    numUGrad + numPGrad + numExchange + numAudit  
from   Courses natural join Profs natural join Enrollment;
```

Views: Providing Logical Data Independence



- **Logical Schema** - logical structure of data in DBMS
- **Physical Schema** - how the data described by logical schema is physically organized in DBMS
- **External Schema** - A customized view of logical schema
- **Logical (Physical) Data independence**: Insulate users/applications from changes to logical (physical) schema

Conditional Expressions: CASE

Scores

name	marks
Alice	92
Bob	63
Carol	58
Dave	47

name	grade
Alice	A
Bob	B
Carol	C
Dave	D

```
select name, case  
    when marks >= 70 then 'A'  
    when marks >= 60 then 'B'  
    when marks >= 50 then 'C'  
    else 'D'  
end as grade  
from Scores;
```

Conditional Expressions: CASE (cont.)

```
case  
  when condition1 then result1  
  ...  
  when conditionn then resultn  
  else result0  
end
```

```
case expression  
  when value1 then result1  
  ...  
  when valuen then resultn  
  else result0  
end
```

Conditional Expressions: COALESCE

Tests

name	first	second	third
Alice	pass	null	null
Bob	fail	pass	null
Carol	fail	fail	pass
Dave	fail	fail	fail
Eve	fail	fail	null

name	result
Alice	pass
Bob	pass
Carol	pass
Dave	fail
Eve	fail

```
select name, case  
    when (first = 'pass') or (second = 'pass')  
        or (third = 'pass') then 'pass'  
    else 'fail'  
end as result  
from Tests;
```

Conditional Expressions: COALESCE

(cont.)

Tests

name	first	second	third
Alice	pass	null	null
Bob	fail	pass	null
Carol	fail	fail	pass
Dave	fail	fail	fail
Eve	fail	fail	null

name	result
Alice	pass
Bob	pass
Carol	pass
Dave	fail
Eve	fail

select name, **coalesce**(third,second,first) **as** result
from Tests;

- **coalesce** returns the first non-null value in its arguments
- Returns null if all the arguments are null

Conditional Expressions: NULLIF

Tests

name	result
Alice	absent
Bob	fail
Carol	pass
Dave	absent
Eve	pass

name	status
Alice	null
Bob	fail
Carol	pass
Dave	null
Eve	pass

select name, **nullif**(result,'absent') as status
from Tests;

- **nullif** ($value_1$, $value_2$)
- Returns *null* if $value_1$ is equal to $value_2$; otherwise returns $value_1$

Pattern Matching with LIKE Operator

Find customer names ending with “e” that consists of at least four characters

select cname **from** Customers **where** cname **like** '___%e';

Customers

cname	area
Homer	West
Lisa	South
Maggie	East
Moe	Central
Ralph	Central
Willie	North

cname
Maggie
Willie

- Underscore **_** matches any single character
- Percent **%** matches any sequence of 0 or more characters
- “string **not like** pattern” is equivalent to “**not** (string **like** pattern)”
- For more advanced regular expressions, use **similar to** operator

Queries with Universal Quantification

- **Example:** Find the names of all students who have enrolled in **all** the courses offered by CS department

Courses (courseId, name, dept)

Students (studentId, name, birthDate)

Enrolls (sid, cid, grade)

Queries with Universal Quantification (cont.)

- Let R denote the set of all students who have enrolled in **all** the courses offered by CS department
- Let $\overline{R} = \text{Students} - R$
- \overline{R} = the set of all students who have **not** enrolled in all the courses offered by CS department
- A student $s \in \overline{R}$ iff there exists some CS course c such that s has not enrolled in c
- Given a studentId x , let $F(x)$ = set of courseIds of CS courses that are not enrolled by student with studentId x
- $\overline{R} = \{s \in \text{Students} \mid F(s.\text{studentId}) \neq \emptyset\}$

Queries with Universal Quantification (cont.)

- $\overline{R} = \{s \in \text{Students} \mid F(s.\text{studentId}) \neq \emptyset\}$
- \overline{R} can be computed by the following pseudo SQL query:

select s.studentId **from** Students **where exists** (F(s.studentId))

- R can be computed by the following pseudo SQL query:

select s.studentId **from** Students **where not exists**
(F(s.studentId))

Queries with Universal Quantification (cont.)

--F(x): set of courseIds of CS courses that are not enrolled
--by student with studentId x

```
select courseId
from   Courses C
where dept = 'CS'
and    not exists (
    select 1
    from   Enrolls E
    where E.cid = C.courseId
    and    E.sid = x
);
```

Queries with Universal Quantification (cont.)

--Names of students who have enrolled in all CS Courses

```
select name
from   Students S
where not exists (
    select courseId
    from   Courses C
    where dept = 'CS'
    and not exists (
        select 1
        from   Enrolls E
        where E.cid = C.courseId
        and   E.sid = S.studentId
    )
);
```

Summary

- Conceptual evaluation of queries

select	distinct select-list
from	from-list
where	where-condition
group by	groupby-list
having	having-condition
order by	orderby-list
limit	limit-specification
offset	offset-specification

- Non-scalar subqueries can be used in FROM, WHERE, and HAVING clauses
- Aggregate functions can be used in SELECT, HAVING, and ORDER BY clauses
- SQL Reference: <https://www.postgresql.org/docs/current/index.html>