

Q1. Consider a table **exams**( sid, cid, **score** ), such that

- Each **sid** is an integer and represents a student ID.
- Each **cid** is an integer and represents a course ID.
- Each **score** is an integer and represents a final exam score of a student in a course.

Write a function **max\_min** with the following properties:

- It has an input parameter **stu\_id**, which is an integer.
- It has two output parameters **max\_cid** and **min\_cid**, both of which are integers.
- It examines the records in **exams** whose **sids** are equal to **stu\_id**, and identifies the two records among them with the largest and smallest **scores**, respectively. (Ties are broken arbitrarily.) For the record with the largest **score**, its **cid** is assigned to **max\_cid**. For the record with the smallest **score**, if its **score** is smaller than the largest **score**, then its **cid** is assigned to **min\_cid**; otherwise, **min\_cid** is set to NULL.

A template for **max\_min** is provided below.

```
CREATE OR REPLACE FUNCTION max_min( IN stu_id integer, OUT max_cid integer, OUT min_cid integer )
RETURNS RECORD as $$
DECLARE
    max_score integer;
    min_score integer;
BEGIN
    // Write your code here
END;
$$ LANGUAGE plpgsql;
```

Q2. Consider the **exams** table in Question Q1. Write a function **revised\_avg** that returns the “revised average score” of a given student, with the following properties:

- The function has an input parameter **stu\_id**, which is an integer.
- The function has one output parameter **r\_avg**, which is a numeric.
- The function examines the records in **exams** whose **sids** are equal to **stu\_id**. If there exist at least 3 such records, the function returns the average score of these records, excluding one record with the highest score (with ties broken arbitrarily) and one record with the lowest score (with ties broken arbitrarily). If there exist less than 3 such records, the function returns NULL.

A template for **revised\_avg** is provided below.

```
CREATE OR REPLACE FUNCTION revised_avg( IN stu_id integer, OUT r_avg float )  
RETURNS float as $$
```

```
// Write your code here
```

```
$$ LANGUAGE plpgsql;
```

Q3. Consider the **exams** table in Question Q1 and the concept of “revised average score” in Question Q2. Write a function **list\_r\_avg** that returns the **sid** of each student in **exams** along with his/her revised average score. For simplicity, we assume that all **sids** in **exams** are non-negative integers.

A template for **list\_r\_avg** is provided below.

```
CREATE OR REPLACE FUNCTION list_r_avg()
RETURNS TABLE ( stu_id integer, ravg float ) AS $$
DECLARE
    curs CURSOR FOR (SELECT sid, score from exams order by sid);
    // write your code here
BEGIN
    // write your code here
END;
$$ LANGUAGE plpgsql;
```

Q4. Consider the **exams** table in Question Q1. Write a function **list\_scnd\_highest** that returns the **sid** of each student in **exams** along with his/her 2nd highest score. (Ties are broken arbitrarily.) If the student has fewer than 2 scores, then **list\_scnd\_highest** returns NULL as his/her 2nd highest score. For simplicity, we assume that all **sids** and **scores** in **exams** are non-negative integers.